



COL333/671: Introduction to AI

Semester I, 2024-25

Local Search Algorithms

Rohan Paul

Outline

- Last Class
 - Informed Search
- This Class
 - Local Search Algorithms
- Reference Material
 - AIMA Ch. 4.1


Acknowledgement

These slides are intended for teaching purposes only. Some material has been used/adapted from web sources and from slides by Doina Precup, Dorsa Sadigh, Percy Liang, Mausam, Dan Klein, Nicholas Roy and others.

Search Methods for Discrete Optimization

- **Setting**

- A set of discrete states, X .
- An objective/evaluation function assigns a “goodness” value to a state, $Eval(X)$
- Problem is to search the state space for the state, X^* that maximizes the objective.



- **Searching for the optimal solution can be challenging. Why?**

- The number of states is very large.
 - Cannot simply enumerate all states and find the optimal.
- We can only evaluate the function.
 - Cannot write it down analytically and optimize it directly.

Key Idea

- Searching for “the optimal” solution is very difficult.
- Question is whether we can search for a reasonably good solution.

Example – Windmill Placements

Problem: Optimizing the locations of windmills in a wind farm

- An area to place windmills.
- Location of windmills affects the others. Reduced efficiency for those in the wake of others.
- Grid the area into bins.
- A large number of configurations of windmills possible.
- Given a configuration we can evaluate the total efficiency of the farm.
- Can neither enumerate all configurations nor optimize the power efficiency function analytically.
- **Goal is to search for the configuration that maximizes the efficiency.**





Figure 5: Turbines experiencing multiple wakes. As an example, turbine 3 is experiencing wake effects from both turbine 1 and 2. Image adopted from [4].

Example: Conference Scheduling

ICRA 2024 Technical Program Tuesday May 14, 2024 [Tuesday](#) [Wednesday](#) [Thursday](#) [Next](#) [Top](#)
Click on the day's program and use the arrow keys for easy horizontal and vertical scrolling

Keynote 1	Award Session 1	Keynote 2	Award Session 2	Keynote 3	Technical Session 1	Technical Session 2	Technical Session 3	Technical Session 4	Technical Session 5	Technical Session 6	Technical Session 7
	10:30-12:00 CC-Main Hall Award Session TuAA1-CC Automation		10:30-12:00 CC-301 Award Session TuAA2-CC Cognitive Robotics		10:30-12:00 CC-303 Oral Session TuAT1-CC Planning under Uncertainty I	10:30-12:00 CC-311 Oral Session TuAT2-CC Mechanism Design I	10:30-12:00 CC-313 Oral Session TuAT3-CC Formal Methods in Robotics and Automation I	10:30-12:00 CC-315 Oral Session TuAT4-CC Multi-Robot Systems I	10:30-12:00 CC-411 Oral Session TuAT5-CC Sensors and Audition	10:30-12:00 CC-414 Oral Session TuAT6-CC 2D/3D Visual Perception	
	13:30-15:00 CC-Main Hall Award Session TuBA1-CC Human-Robot Interaction		13:30-15:00 CC-301 Award Session TuBA2-CC Mechanisms and Design		13:30-15:00 CC-303 Oral Session TuBT1-CC Planning under Uncertainty II	13:30-15:00 CC-311 Oral Session TuBT2-CC Mechanism Design II	13:30-15:00 CC-313 Oral Session TuBT3-CC Formal Methods in Robotics and Automation II	13:30-15:00 CC-315 Oral Session TuBT4-CC Multi-Robot Systems II	13:30-15:00 CC-411 Oral Session TuBT5-CC Vision Systems	13:30-15:00 CC-414 Oral Session TuBT6-CC RGB-D Sensing and Perception I	
15:30-16:00 National Convention Hall Keynote Session TuKN1-CC		15:30-16:00 CC-Main Hall Keynote Session TuKN2-CC		15:30-16:00 CC-301 Keynote Session TuKN3-CC							

ICLR My Stuff [Login](#)

Select Year: (2024) Getting Started Schedule Main Conference Workshops Community Sponsors Organizers Help

Show Detail [Schedule](#) Tue Wed Thu Fri Sat Timezone: [America/Los_Angeles](#)

Filter Events: Nothing selected Filter Rooms: Nothing selected


MON 6 MAY	TUE 7 MAY	WED 8 MAY	THU 9 MAY	FRI 10 MAY	SAT 11 MAY
10 p.m. Registration Desk (ends 10:00 AM)	12:30 a.m. Break: Coffee Break (ends 1:00 AM)	12:30 a.m. Break: Coffee Break (ends 1:00 AM)	12:30 a.m. Break: Coffee Break (ends 1:00 AM)	12:30 a.m. Break: Coffee Break (ends 1:00 AM)	midnight Workshop: Machine Learning for Genomics Explorations (MLGenX) (ends 8:00 AM)
11:15 p.m. Remarks: Opening Remark (ends 11:30 PM)	1 a.m. Oral 1A <ul style="list-style-type: none">Predictive auxiliary objectives in deep RL mimic learning in the brainClimODE: Climate and Weather Forecasting with Physics-informed Neural ODEsProtein Discovery with Discrete Walk-Jump Sampling	1 a.m. Oral 3A <ul style="list-style-type: none">LoftQ: LoRA-Fine-Tuning-aware Quantization for Large Language ModelsPhenomenal Yet Puzzling: Testing Inductive Reasoning Capabilities of Language Models with Hypothesis RefinementReLU Strikes Back: Exploiting Activation Sparsity in Large Language Models	1 a.m. Oral 5A <ul style="list-style-type: none">Generalization in diffusion models arises from geometry-adaptive harmonic representationsAn Analytical Solution to Gauss-Newton Loss for Direct Image AlignmentGenerative Modeling with Phase Stochastic Bridge	1 a.m. Oral 7A <ul style="list-style-type: none">Small-scale proxies for large-scale Transformer training instabilitiesAn Analytical Solution to Gauss-Newton Loss for Direct Image AlignmentStatistically Optimal K-means Clustering via Nonnegative Low-rank Semidefinite Programming	Workshop: Bridging the Gap Between Practice and Theory in Deep Learning (ends 8:00 AM)
Remarks: Opening Remark (ends 11:30 PM)	Oral 1B <ul style="list-style-type: none">A Real-World WebAgent with Planning, Long Context Understanding, and Reasoning Feedback	Oral 3B <ul style="list-style-type: none">Unified Generative Models of 3D Motions	Oral 5B <ul style="list-style-type: none">Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To	Oral 7B <ul style="list-style-type: none">DreamGaussian: Generative Gaussian Splatting for Efficient 3D Content Creation	Workshop: Global AI Cultures (ends 8:00 AM)
Remarks: Opening Remark (ends 11:30 PM)					Workshop: Secure and Trustworthy Large Language Models (ends 8:00 AM)

Assign papers that are similar in a session. Avoid conflicts between sessions.

Example

4-Queens Problem

- Discrete set of states: 4 queens in 4 columns ($4^4 = 256$ states)
- Goal is to find a configuration such that there are no attacks.
 - Moving a piece will change the configuration.
- Any configuration can be evaluated using a function
 - $h(x)$ = number of attacks (number of violated binary constraints)
- Search for the configuration that is optimal such that $h = 0$.



Example

Formally

Variables: x_0, x_1, x_2, x_3 where x_i is the row position of the queen in column i , where $i \in \{0, 1, 2, 3\}$. Assume that there is one queen per column.

Domains: $x_i \in \{0, 1, 2, 3\} \forall i$.

Initial state: 4 queens on the board in random row positions.


Goal state: 4 queens on the board with no pair of queens attacking each other.

Neighbour relation:

- Version A: move a single queen to a different row in the same column.
- Version B: swap the row positions of two queens.

Cost function: the number of pairs of queens attacking each other, directly or indirectly.

Number of attacks are 4.



Local Search Methods

- **Keep track of a single "current" state**
 - We need a principled way to search/explore the state space hoping to find the state with the optimal evaluation.
 - Do not maintain a search tree as we need the solution not the path that led to the solution.
 - Only maintain a single current state.
- **Perform local improvements**
 - Look for alternatives in the vicinity of that solution
 - Try to move towards more better solutions.

Hill-climbing Search



Hill climbing

Let S be the start node and let G be the goal node.

Let $h(c)$ be a heuristic function giving the value of a node

Let c be the start node

Loop

Let $c' =$ the highest valued neighbor of c

If $h(c) \geq h(c')$ then return c

$c = c'$

Start at a configuration. Evaluate the neighbors. Move to the highest valued neighbor if its value is higher than the current state. Else stay.

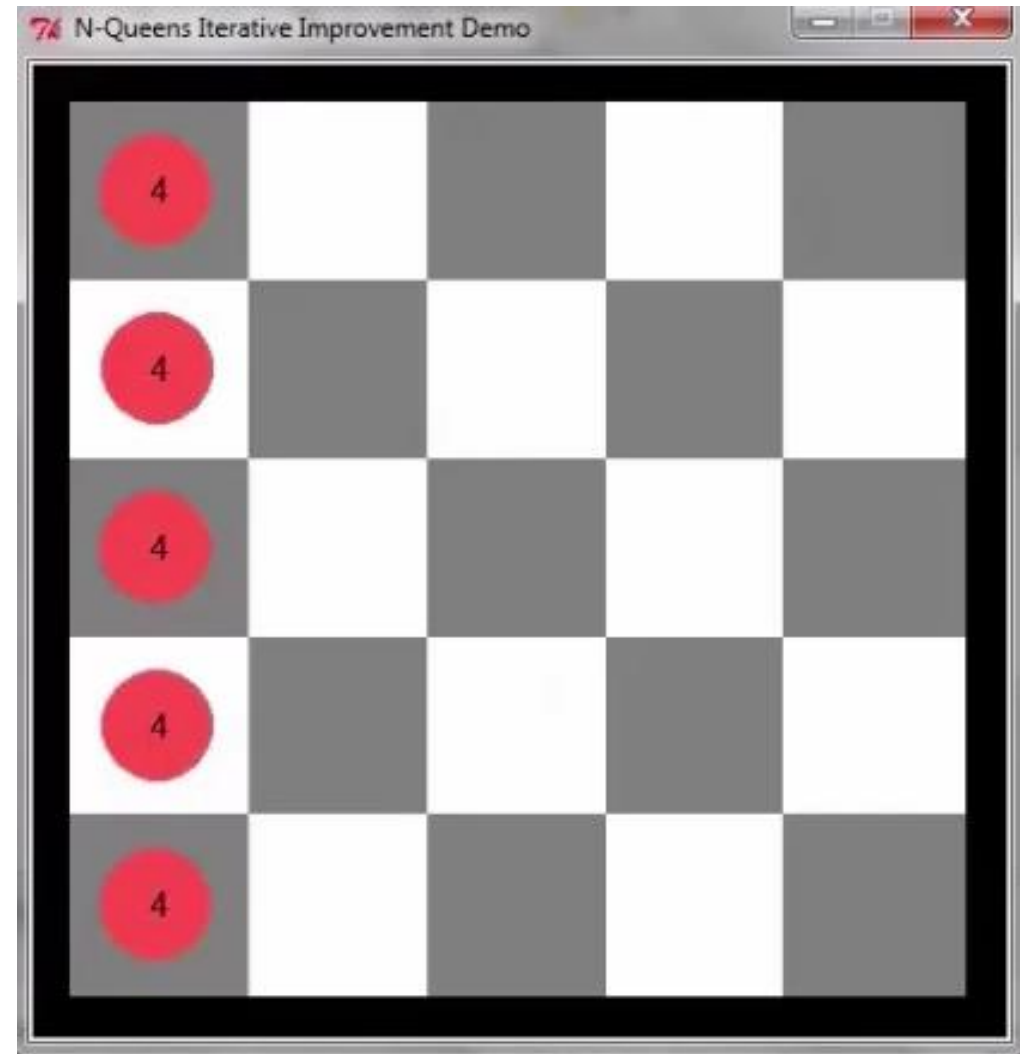
“Like climbing Everest in thick fog with amnesia”

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
  end
```

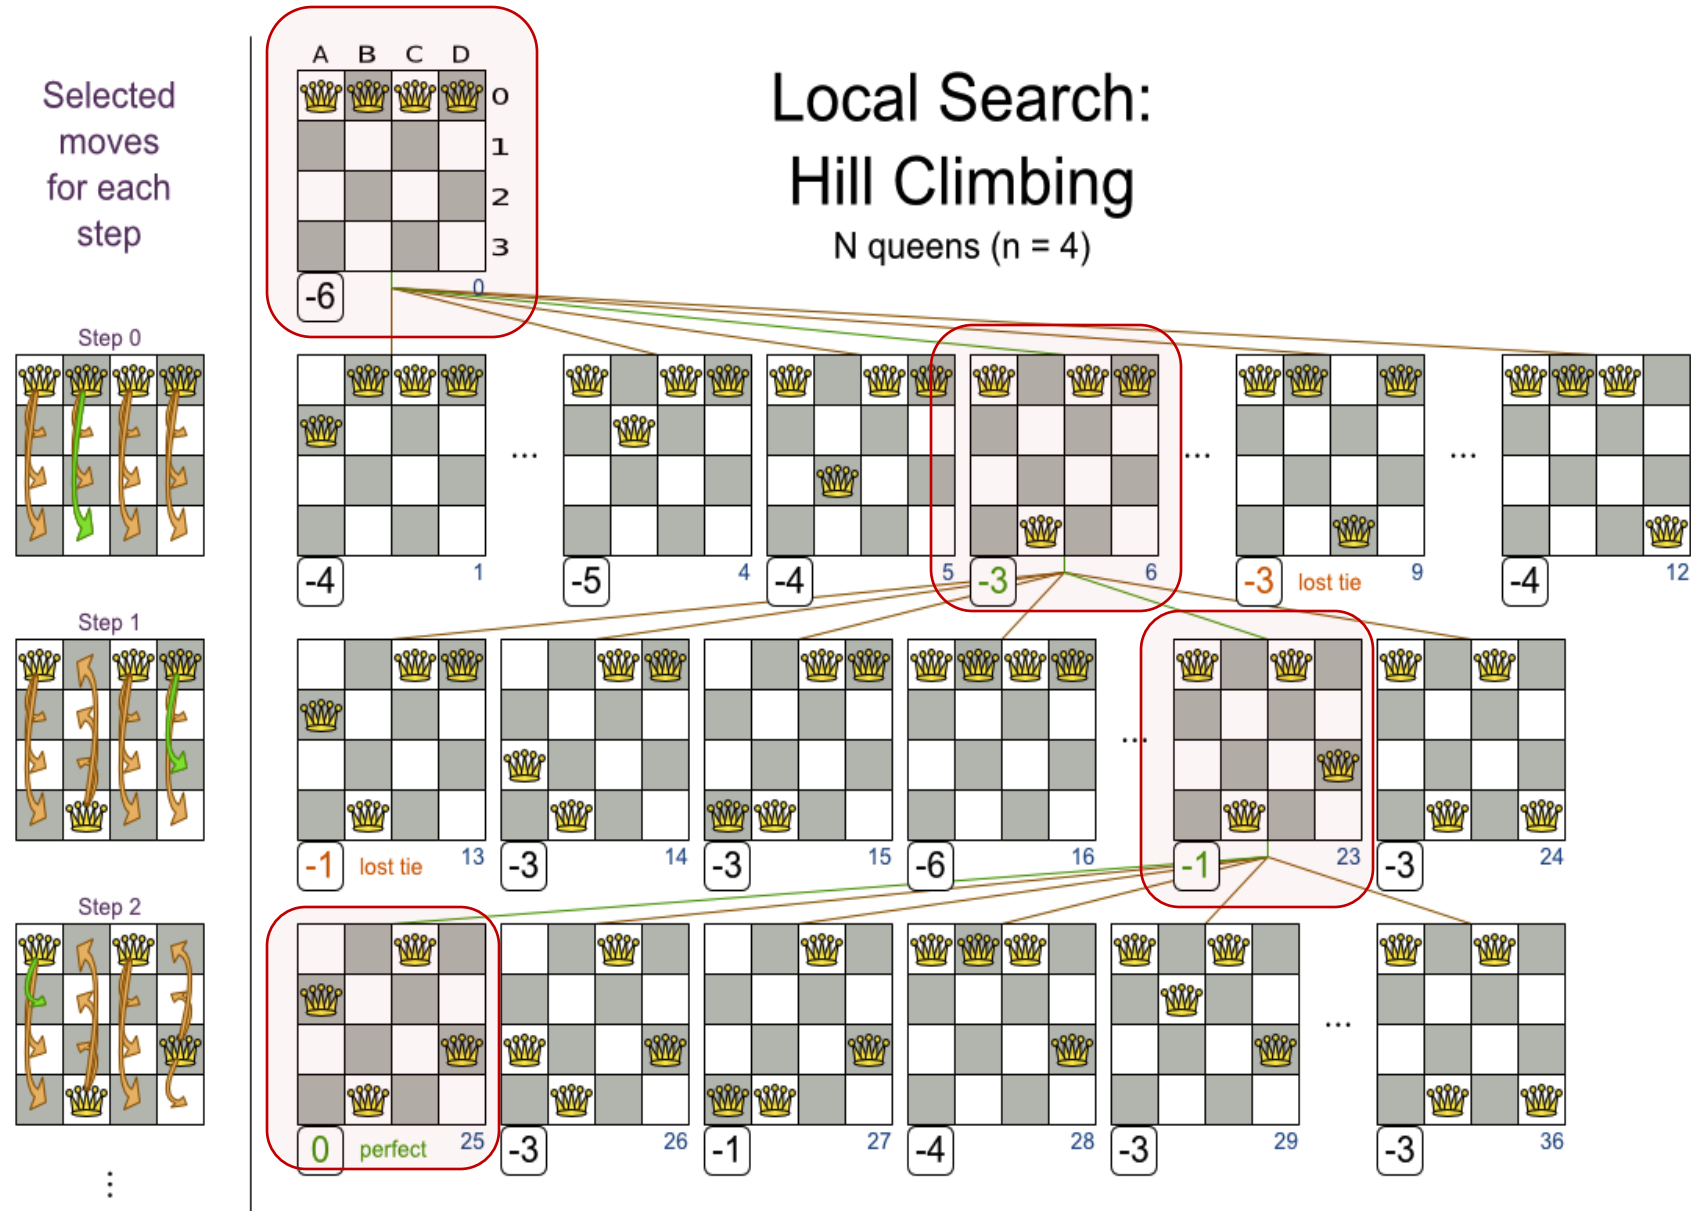
Hill climbing for 4 -queens

- Select a column and move the queen to the square with the fewest conflicts.
- Perform local modifications to the state by changing the position of one piece till the evaluation is minimum.
- Evaluate the possibilities from a state and then jump to that state.



Example


- Local search looks at a state and its local neighborhood.
- Not constructing the entire search tree.
- Consider local modifications to the state. Immediately jump to the next promising neighbor state. Then start again.
- Highly scalable.*



Example: Idea of local improvements

Locally improving a solution for a Travelling Salesperson Problem.




Start with any complete tour, perform pairwise exchanges




Variants of this approach get within 1% of optimal very quickly with thousands of cities

The idea of making local improvements to a candidate solution is a general and widely applicable technique.

8-Queens Problem




Is this an optimal state?



Issue: search reaches a solution where it cannot be improved - a local minimum.

Local minima ($h = 1$). Every successor has a higher cost.

Core Problem in Local Search



- Hill climbing prone to local maxima. Neighbors may not be of higher value. Search will stop at a sub-optimal solution
- Locally optimal actions may not lead to the globally optimal solution

Escaping local minima: Adding randomness


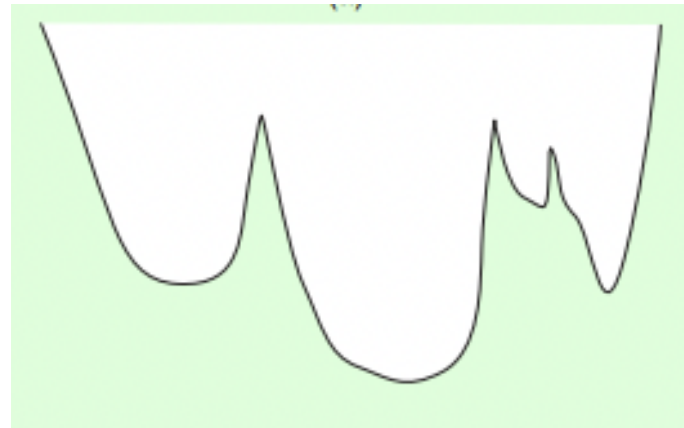
- **Random Re-starts**

- A series of searches from randomly generated initial states.

- **Random Walk**

- Pick "any" candidate move (whether improves the solution or not).

Q: Which method to use for the following cost surfaces?
Random re-starts or random walk?




Escaping local minima: Adding randomness

- **Escaping flat local minima (shoulders)**

- When local search reaches a flat area, that is, when all the neighbours have the same cost as the current state, it terminates right away
- Keep moving - strategy
 - Make sideways moves for a few steps.

- ***Stochastic* Hill Climbing**

- Instead of picking the *best move*, pick *any* move that produces an *improvement*.



Looking for Solution from Multiple Points

- **Local Beam Search**

- Algorithm


- Track k states (rather than 1).
- Begin with k randomly sampled states.
- Loop
 - Generate successors of each of the k-states
 - If anyone has the goal, the algorithm halts
 - Otherwise, select only the k-best successors from the list and repeat.

- Note:

- Each run is not independent, information is passed between parallel search threads.
- Promising states are propagated. Less promising states are not propagated.
- Problem: states become concentrated in a small region of space.

Beam Search is a General Search Technique


- Beam search is a general idea (see right figure).
- Instead of considering all solutions at a level, consider only the top-k.
- Note: usually our memory is finite in size, there is an upper bound on the number of states that can be kept.
- In general, it is an approximate search method.



Beam search is a general idea. Here, shown in the context of a tree search. Beam size is 3. For local search we don't construct the full tree.

“Stochastic” Beam Search

- **Local beam search**
 - Problem: states become concentrated in a small region of space
 - Search degenerates to hill climbing
- **Stochastic beam search**
 - Instead of taking the best k states
 - Sample k states from a distribution
 - Probability of selecting a state *increases* as the *value* of the state.



Instead of top k, sample k given a probability.

Simulated Annealing

- In case of an improving move – move there.
- But allow some apparently **bad moves** - to escape local maxima.
- **Decrease** the size and the frequency of bad moves over time.

- Algorithm sketch



1. Start at initial configuration X of value E (high is good)
2. Repeat:
 - (a) Let X_i be a *random neighbor* of X and E_i be its value
 - (b) If $E < E_i$ then let $X \leftarrow X_i$ and $E \leftarrow E_i$
 - (c) Else, *with some probability p* , still accept the move: $X \leftarrow X_i$ and $E \leftarrow E_i$

- Best solution ever found is always remembered

A form of Monte-Carlo Search. Move around the environment to explore it instead of systematically sweeping. Powerful technique for large domains.

Simulated Annealing: How to decide p ?

- Considering a move from state of value E to a lower valued state of E' . *That is considering a sub-optimal move (E is higher than E').*
- If $(E - E')$ is large:
 - Likely to be close to a promising maximum.
 - Less inclined to go downhill.
- If $(E - E')$ is small:
 - The closest maximum may be shallow
 - More inclined to go downhill is not as bad.



Simulated Annealing: Selecting Moves

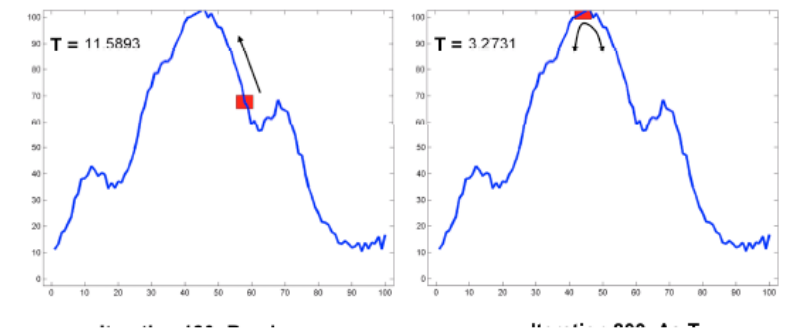


- If the new value E_i is **better** than the old value E , move to X_i
- If the new value is **worse** ($E_i > E$) then move to the neighboring solution as per *Boltzmann* distribution.

$$\exp\left(-\frac{E - E_i}{T}\right)$$

- Temperature ($T > 0$)
 - **T is high**, \exp is ~ 0 , acceptance probability is ~ 1 , high probability of acceptance of a worse solution.
 - **T is low**, the probability of moving to a worse solution is ~ 0 , low probability of acceptance of a worse solution.
 - Schedule T to reduce over time.

Simulated Annealing

- **T is high**
 - The algorithm is in an *exploratory* phase
 - Even bad moves have a high chance of being picked
- **T is low**
 - The algorithm is in an *exploitation* phase
 - The “bad” moves have very low probability
- **If T is decreased slowly enough**
 - Simulated annealing is guaranteed to reach the best solution in the limit.





Able to escape local maxima.

Adding (some) memory: Tabu Search

- Local search loses track of the global cost landscape.
 - May frequently come back to the same state
- Introduce “memory” to prevent re-visits.
 - Maintain a finite-sized “tabu” list which remembers recently visited states so that one does not go towards them.
 - If a state proposed in the neighbourhood is in the tabu list do not go.

Motivating example: PCB layout with lower wire overlaps.



Search with Memory

- Tabu Search
 - Maintain a tabu list of the k last assignments.
 - Don't allow an assignment that is already on the tabu list.
 - If $k = 1$, we don't allow an assignment of to the same value to the variable chosen.
 - Maintain a finite-sized tabu list (a form of local memory) which remembers recently visited states so that one does not go towards them.
 - Note: Tabu search allows for sub-optimal moves.
- Types of memory rules
 - Short-term: immediate states visited in the past.
 - Longer-term: guide the search towards certain regions of the search – where all have we explored in the past.
- Generalise to – searching locally by growing a tree for a short horizon and then picking a move (combining local and tree search).

Genetic Algorithms


- Idea

- Variant of stochastic beam search: progression is by modifying a state.
- Combine two states to generate the successor.
- A mechanism to propose next moves in a different way.

- Ingredients

- Coding of a solution into a string of symbols or bit-string
- A fitness function to judge the worth of a state (or configuration)
- A population of states (or configurations)

= stochastic local beam search + generate successors from **pairs** of states



Many variations:

how selection will be applied, what type of cross-over operators will be used, etc.

Selection of individuals according to a fitness function and pairing


Calculation of the breaking points and recombination

According to a given probability elements in the string are modified.

Genetic Algorithms

https://rednuht.org/genetic_cars_2/

View as a way to propose moves – in an evolutionary way.



Advantage: ability to combine large blocks that evolved independently, impact the granularity of search.